

Guiding unstructured responses

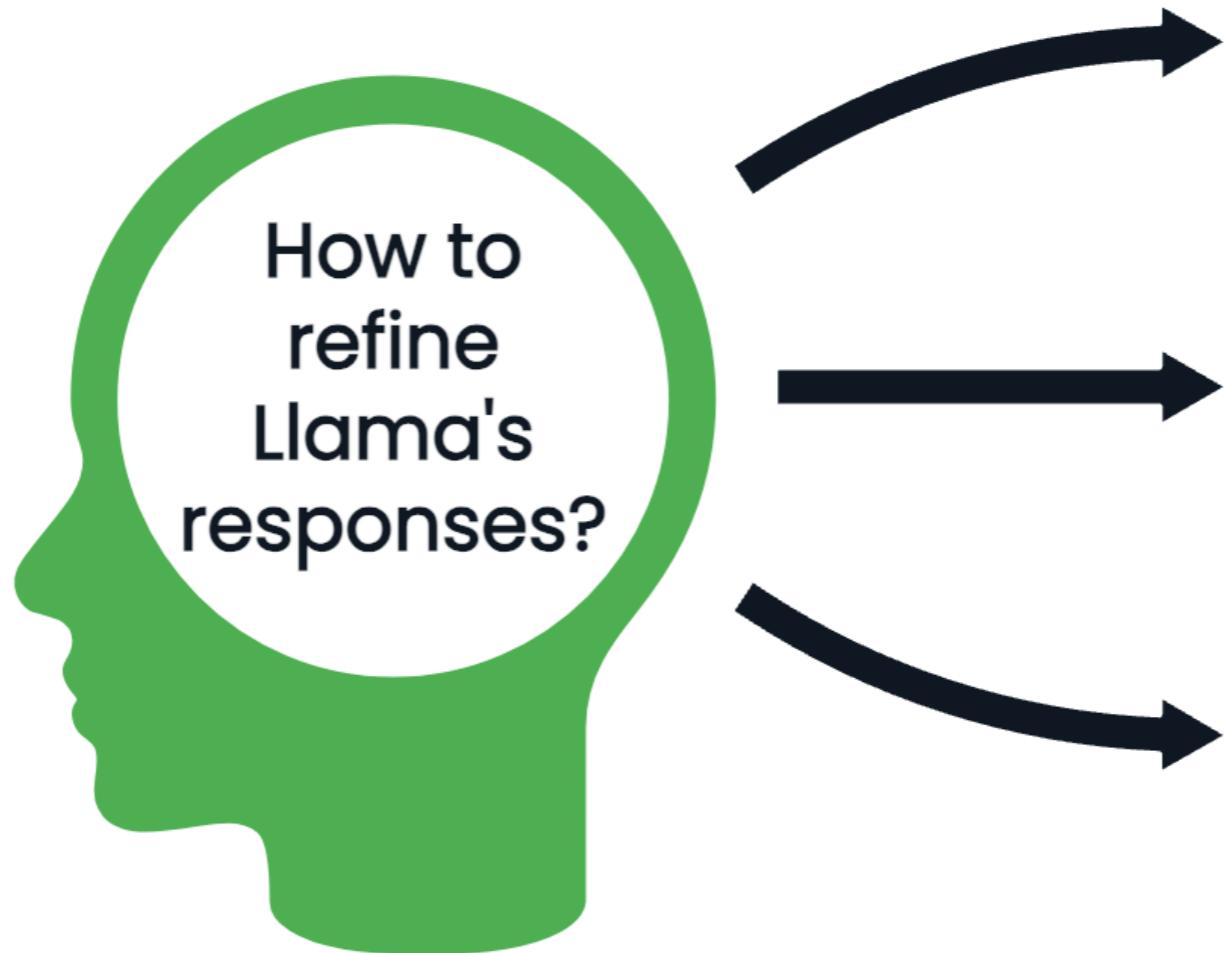
WORKING WITH LLAMA 3



Imtihan Ahmed
Machine Learning Engineer

Controlling model output

- Parameters
- Roles



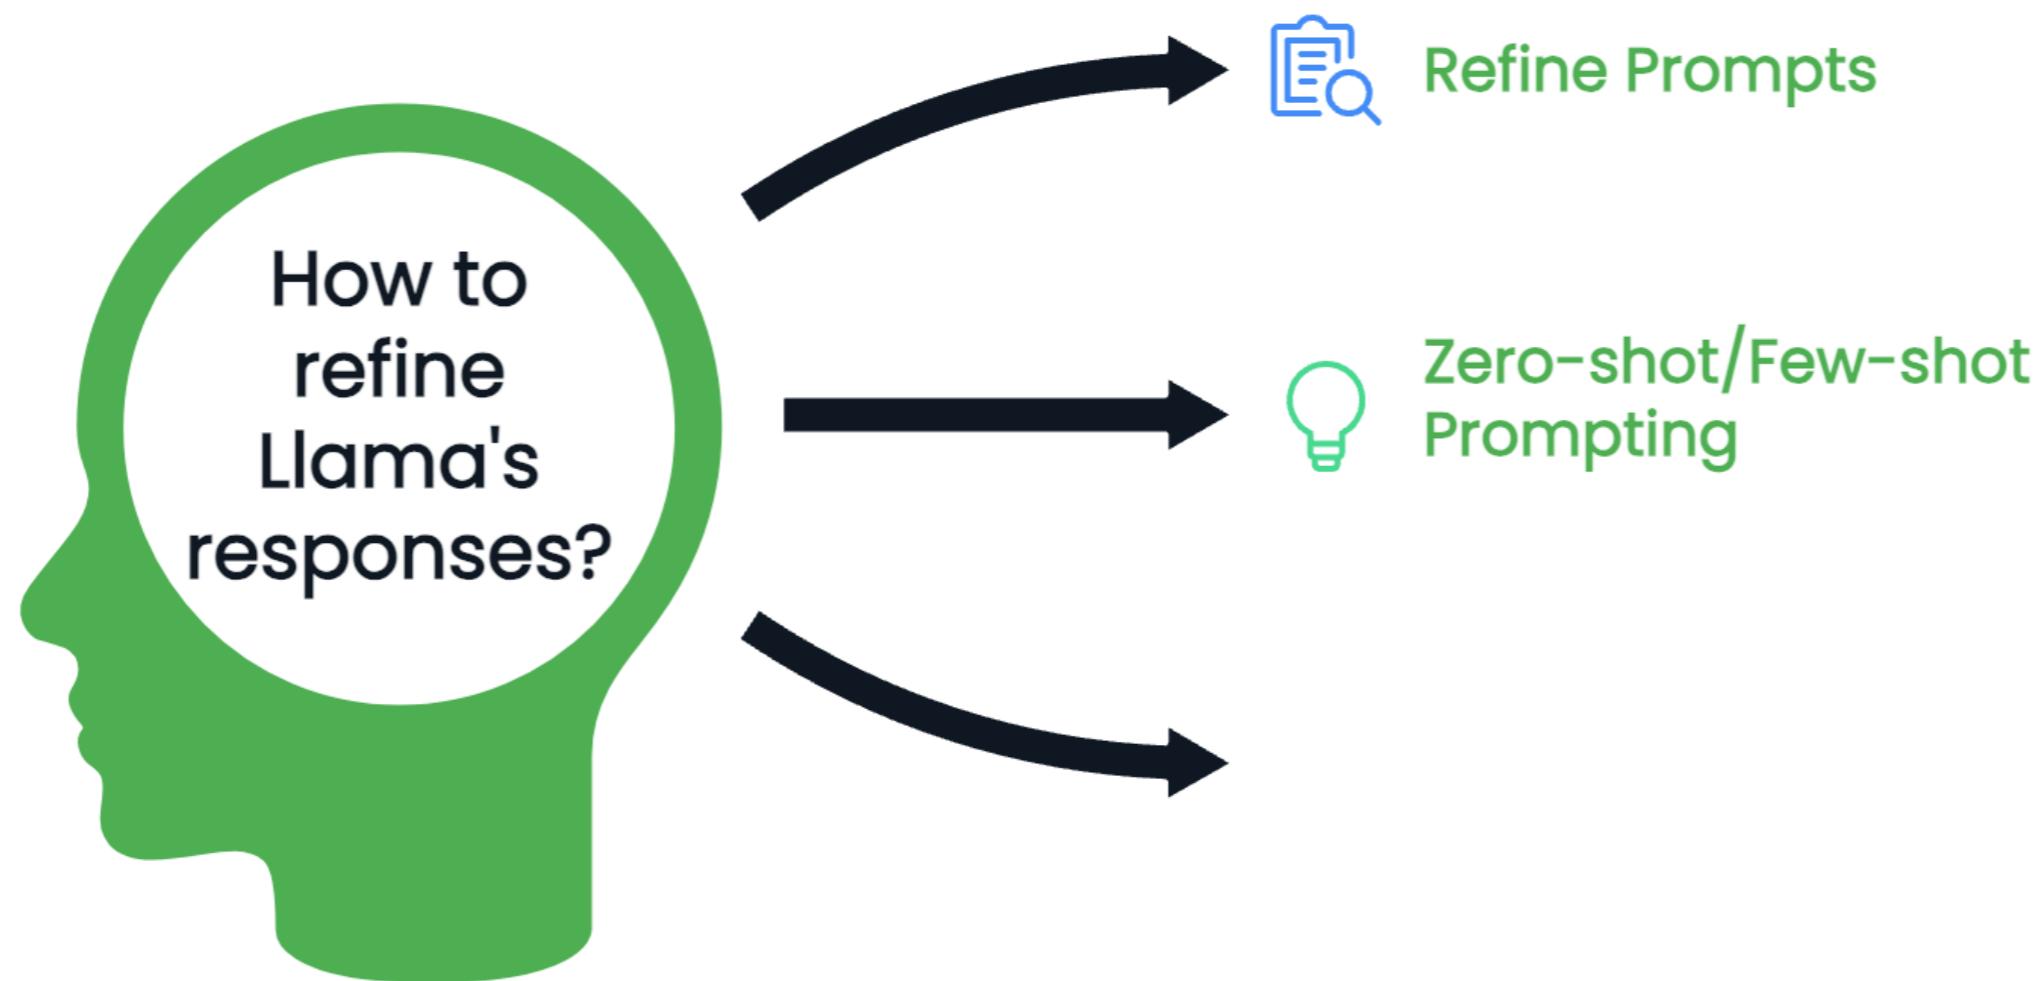
Controlling model output

- Parameters
- Roles



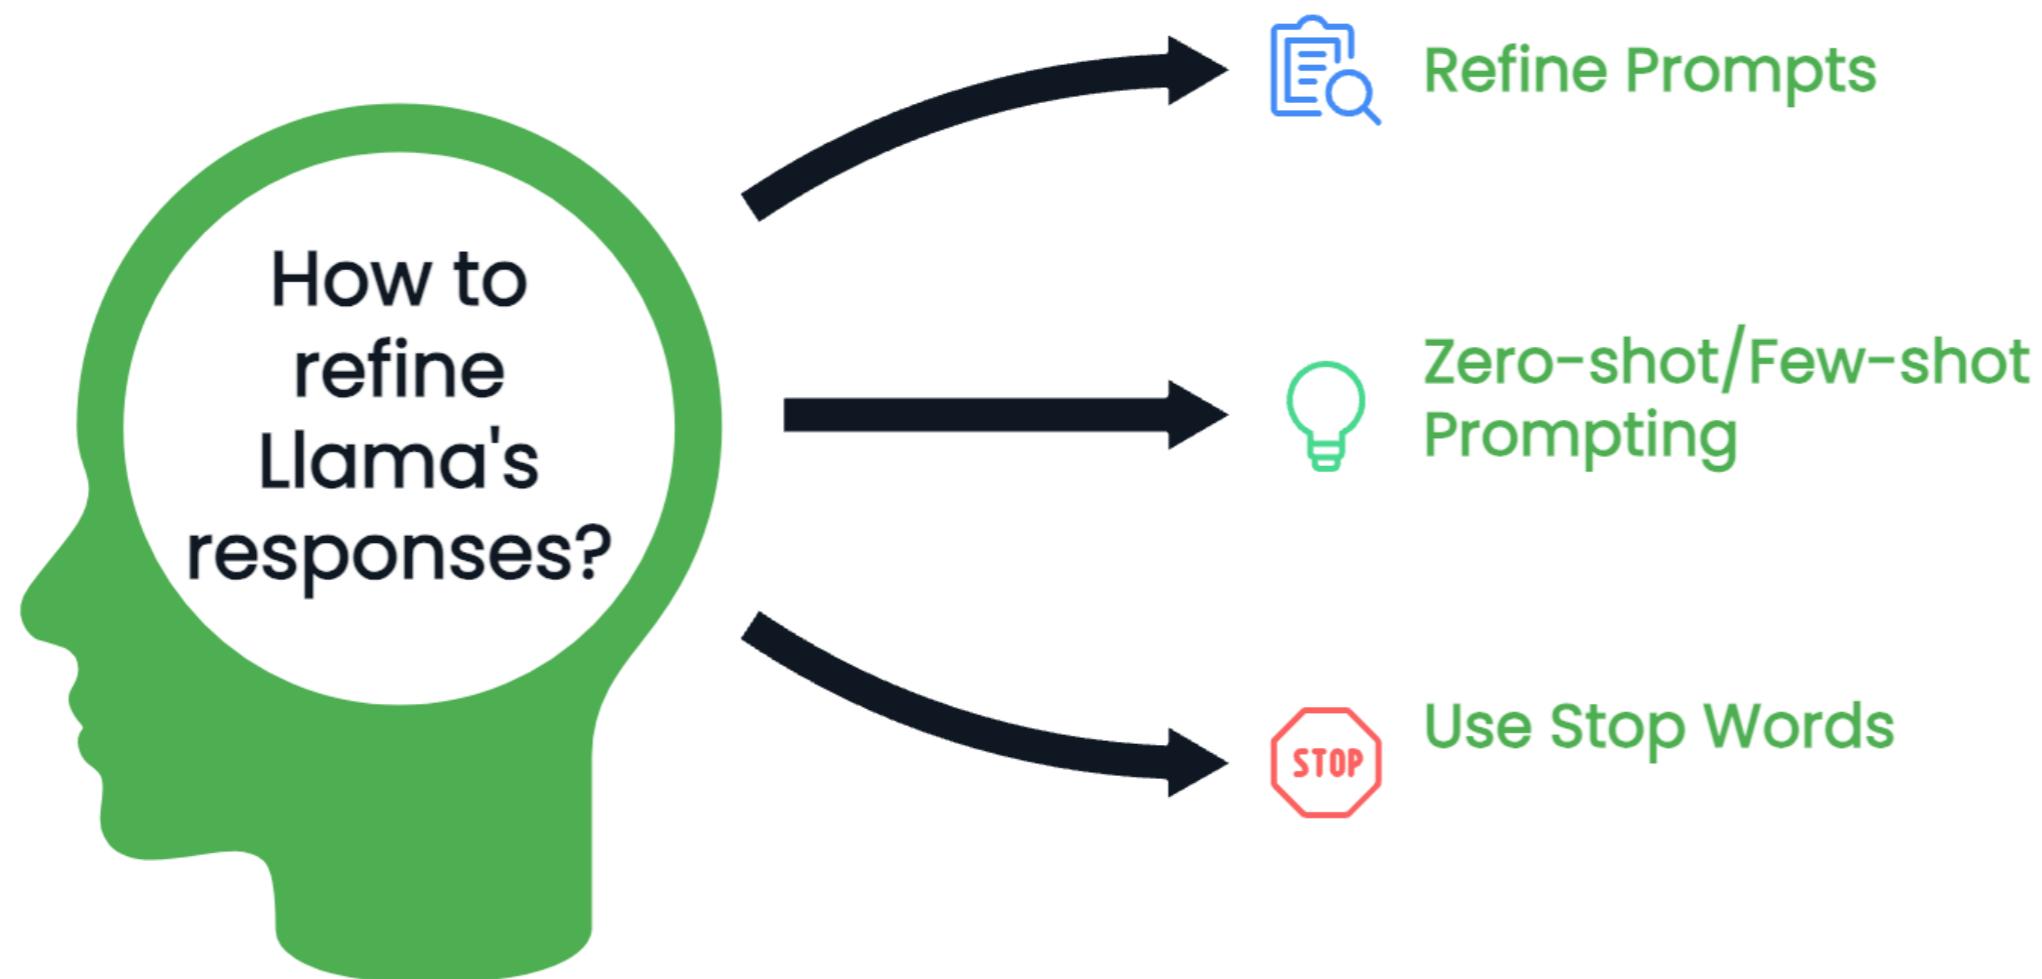
Controlling model output

- Parameters
- Roles



Controlling model output

- Parameters
- Roles



Refining prompts

- Example: summarization

```
text_choice1 = "Summarize key trends in the aviation industry from the last year,  
focusing on fuel efficiency innovations."
```

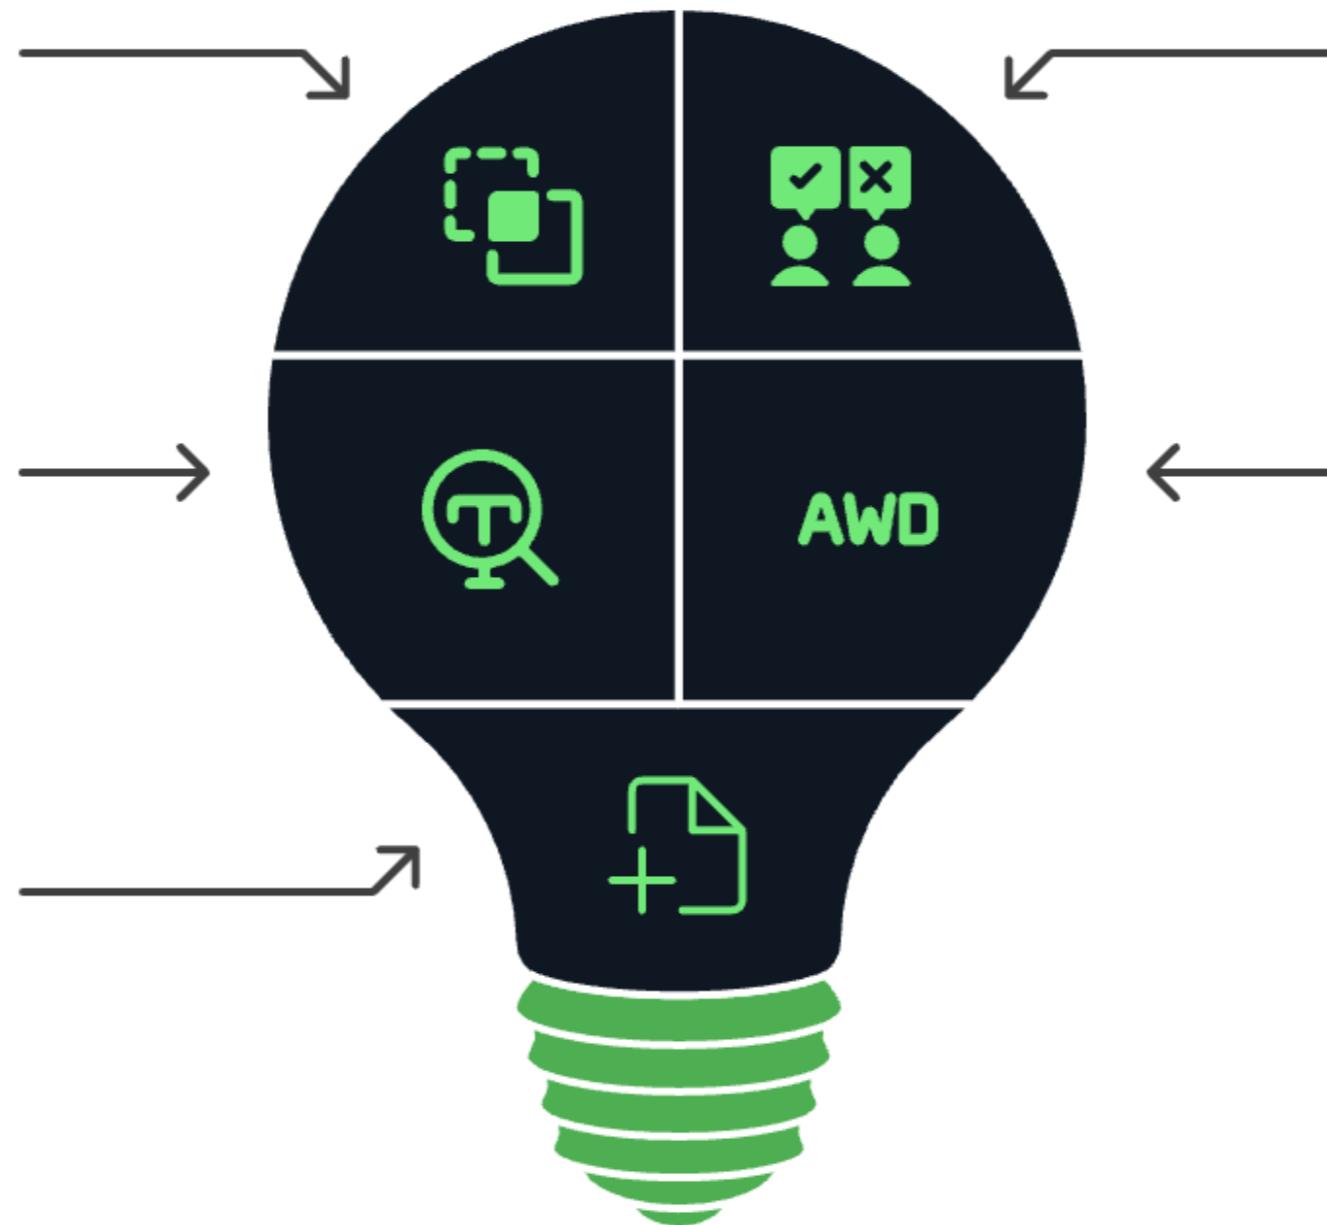
```
text_choice2 = "Tell me about the aviation industry."
```

```
output = llm(text_choice1) # More specific prompt is more effective
```

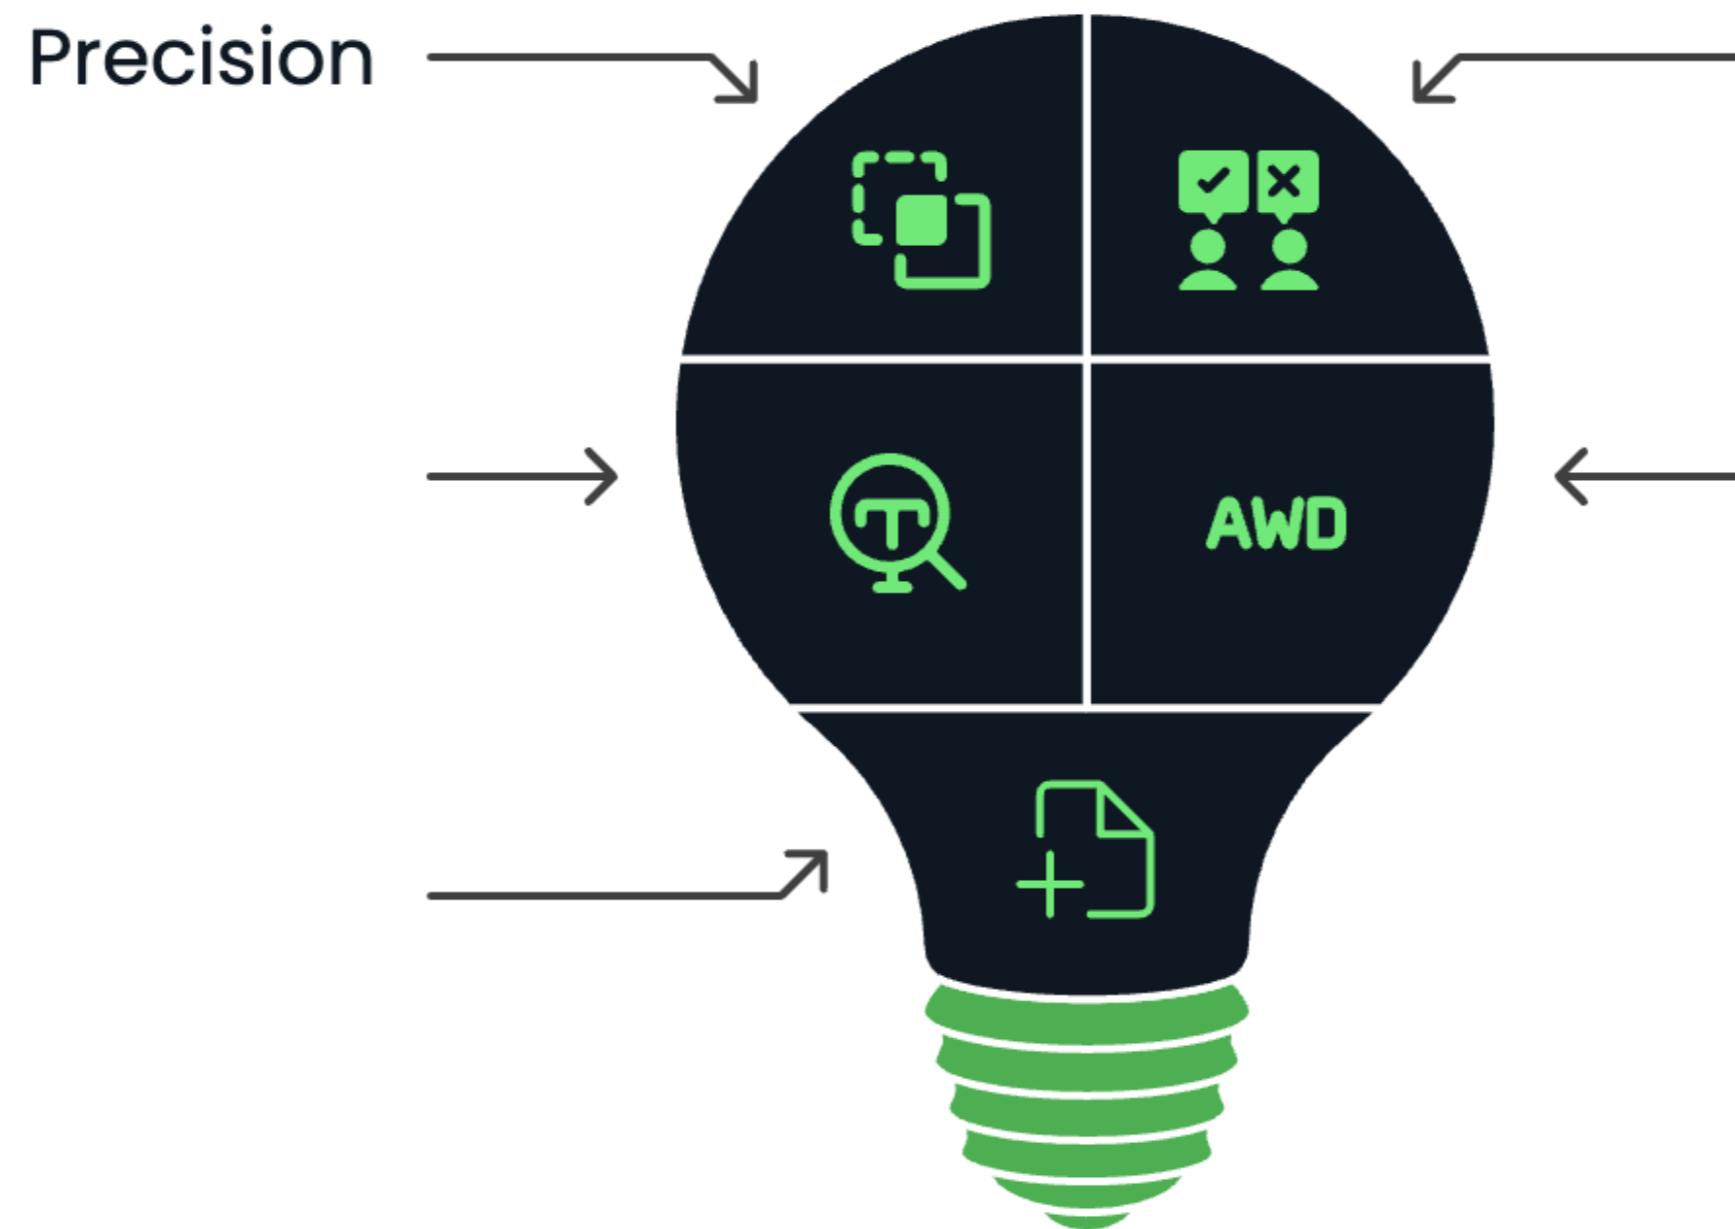
```
print(output['choices'][0]['text'])
```

The aviation industry has made significant strides in fuel efficiency innovations over the last year, driven by the need to reduce greenhouse gas emissions and operating costs. Sustainable Aviation Fuels (SAFs) have emerged as a key trend...

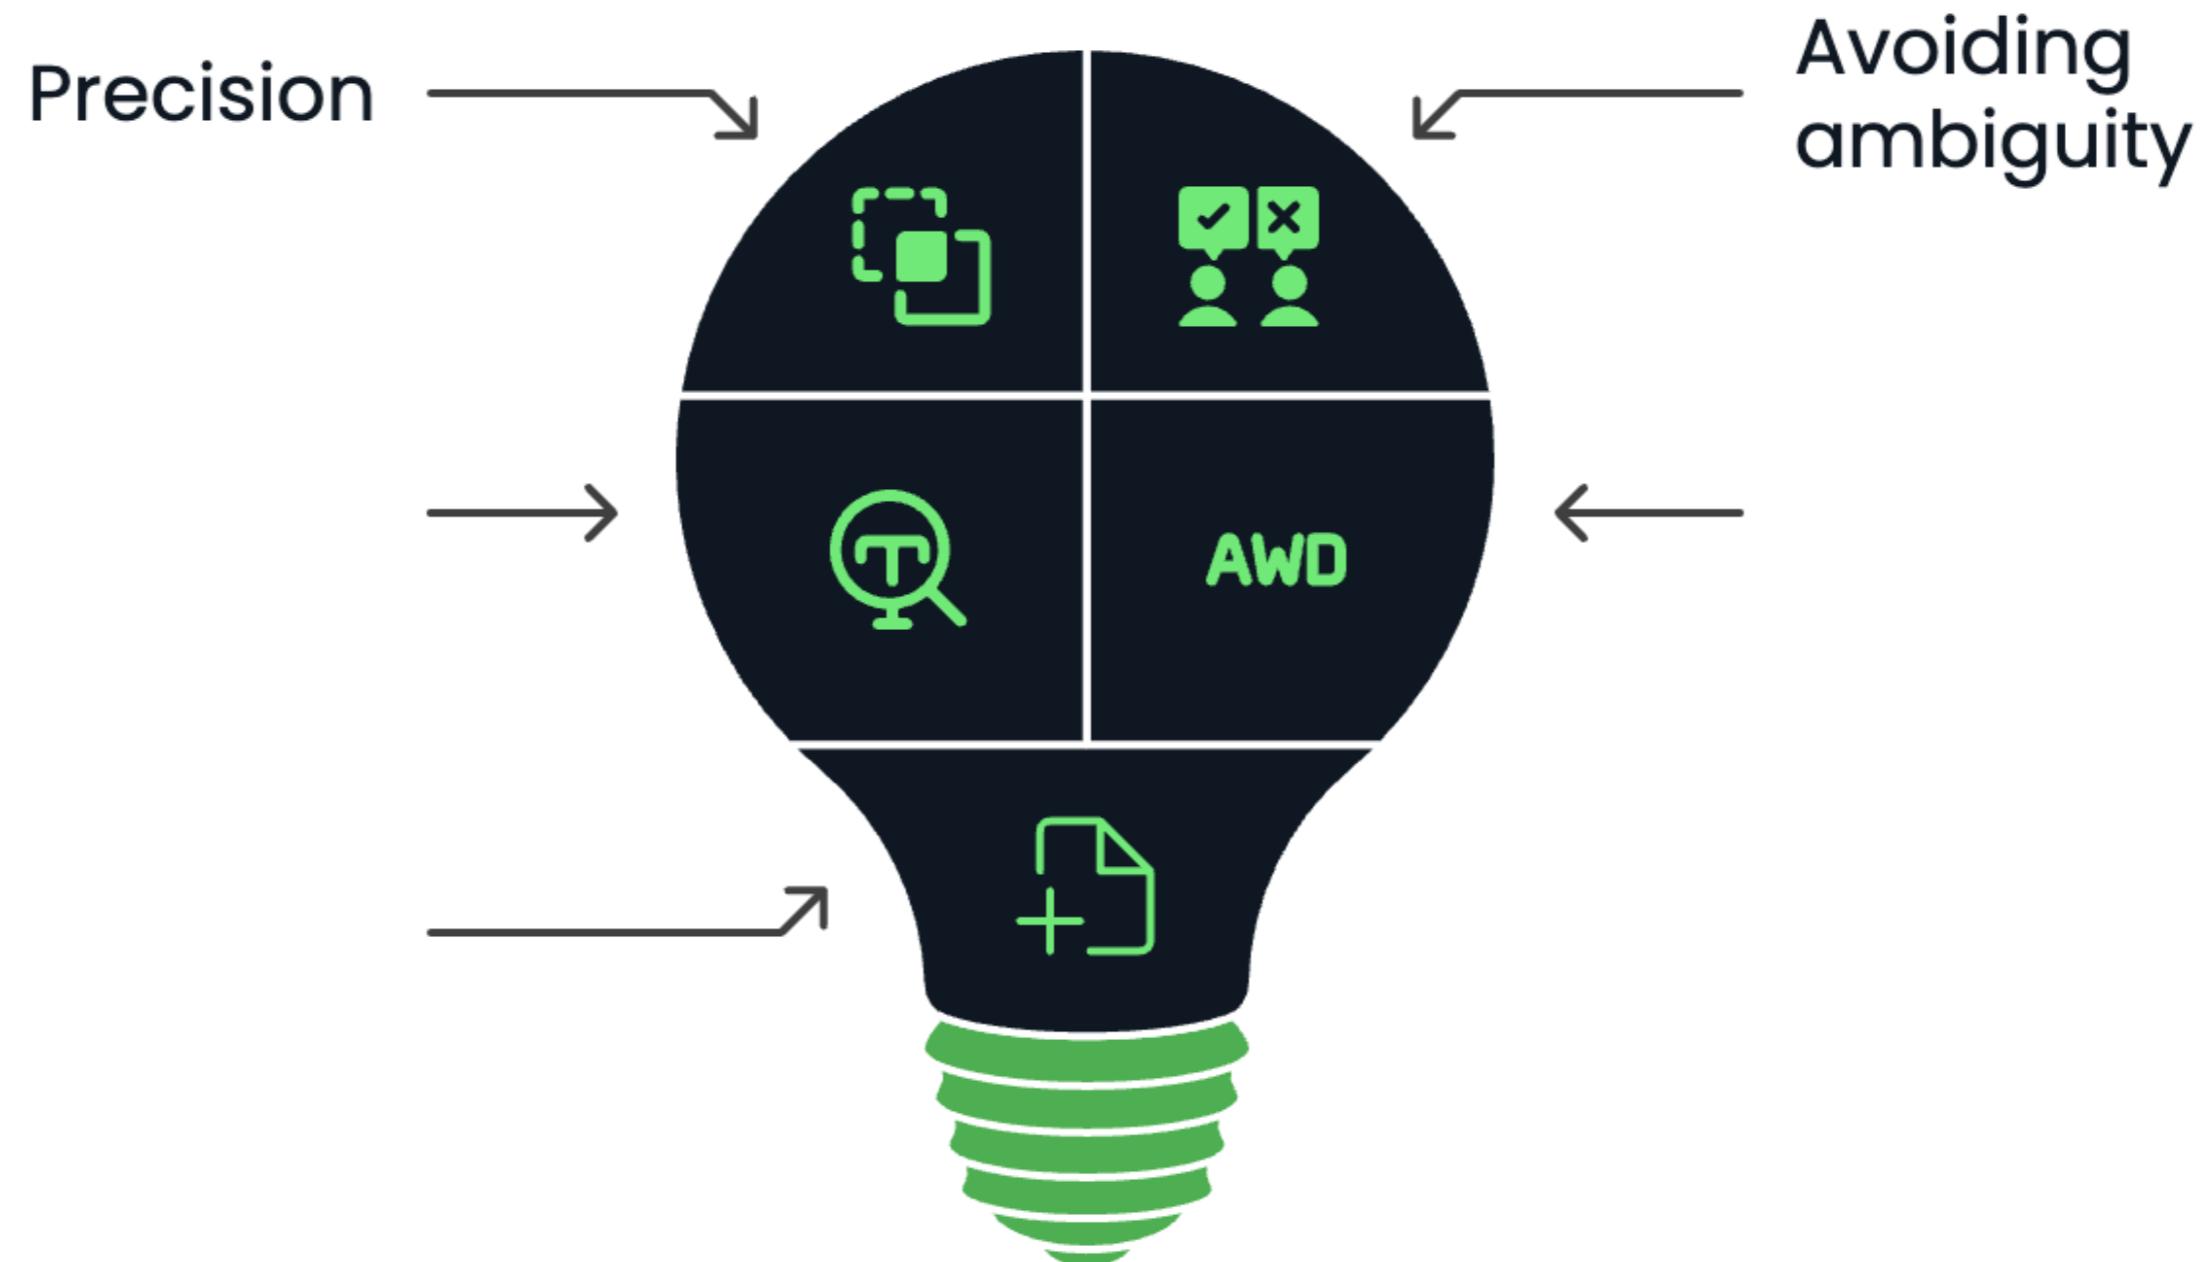
Components of effective prompting



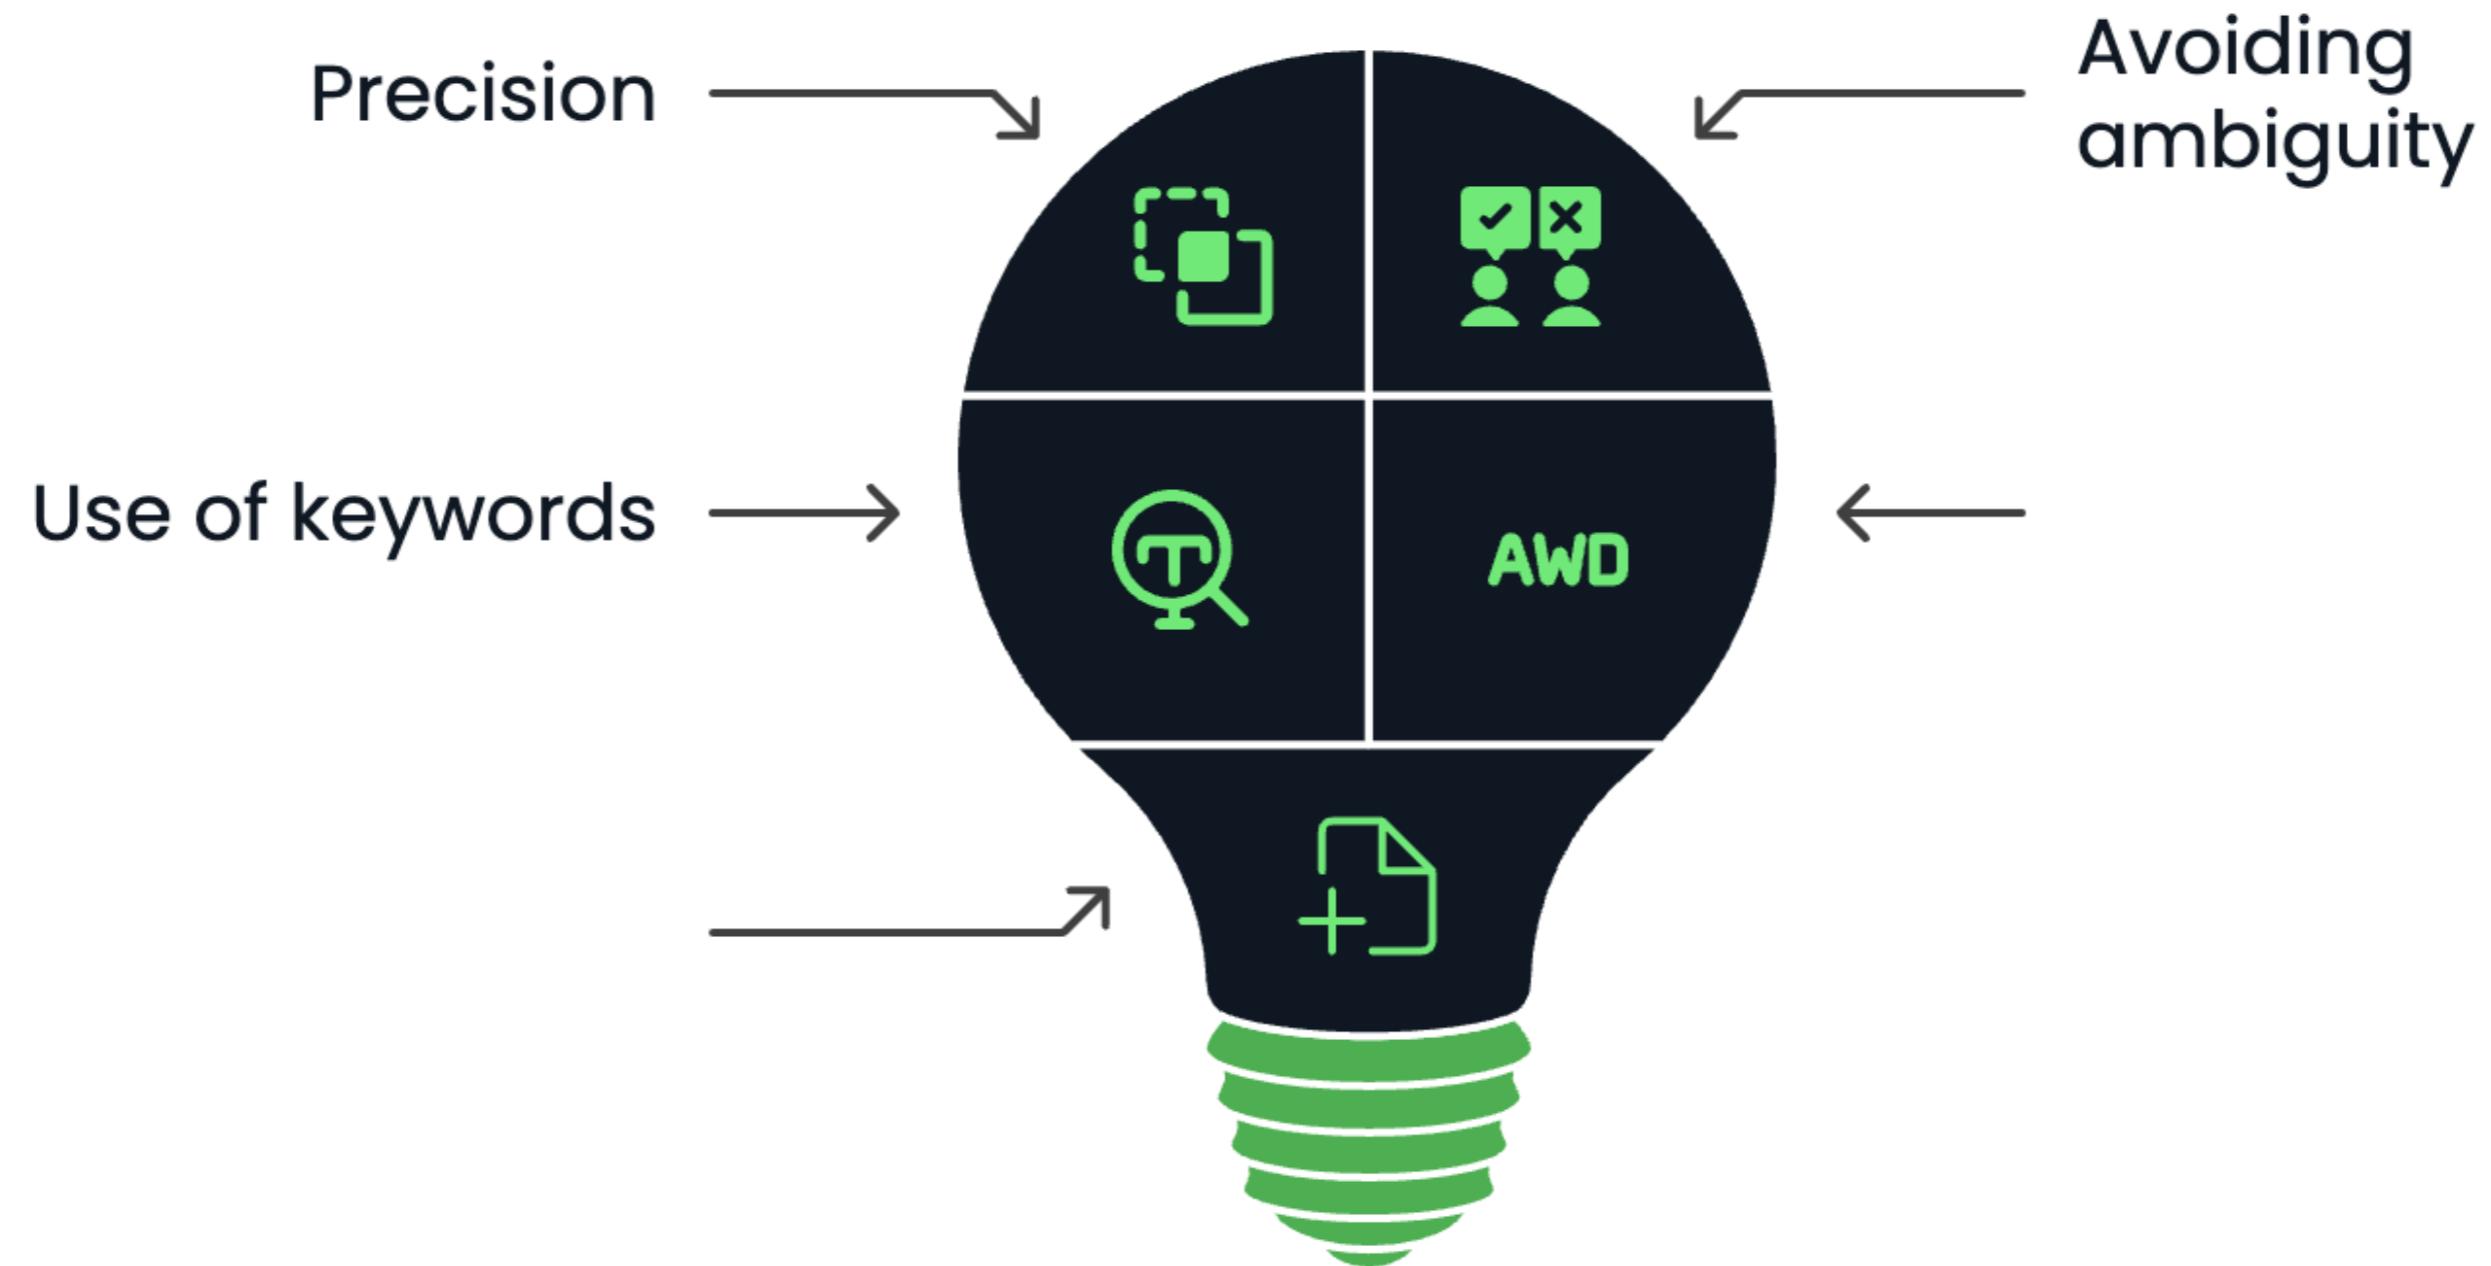
Components of effective prompting



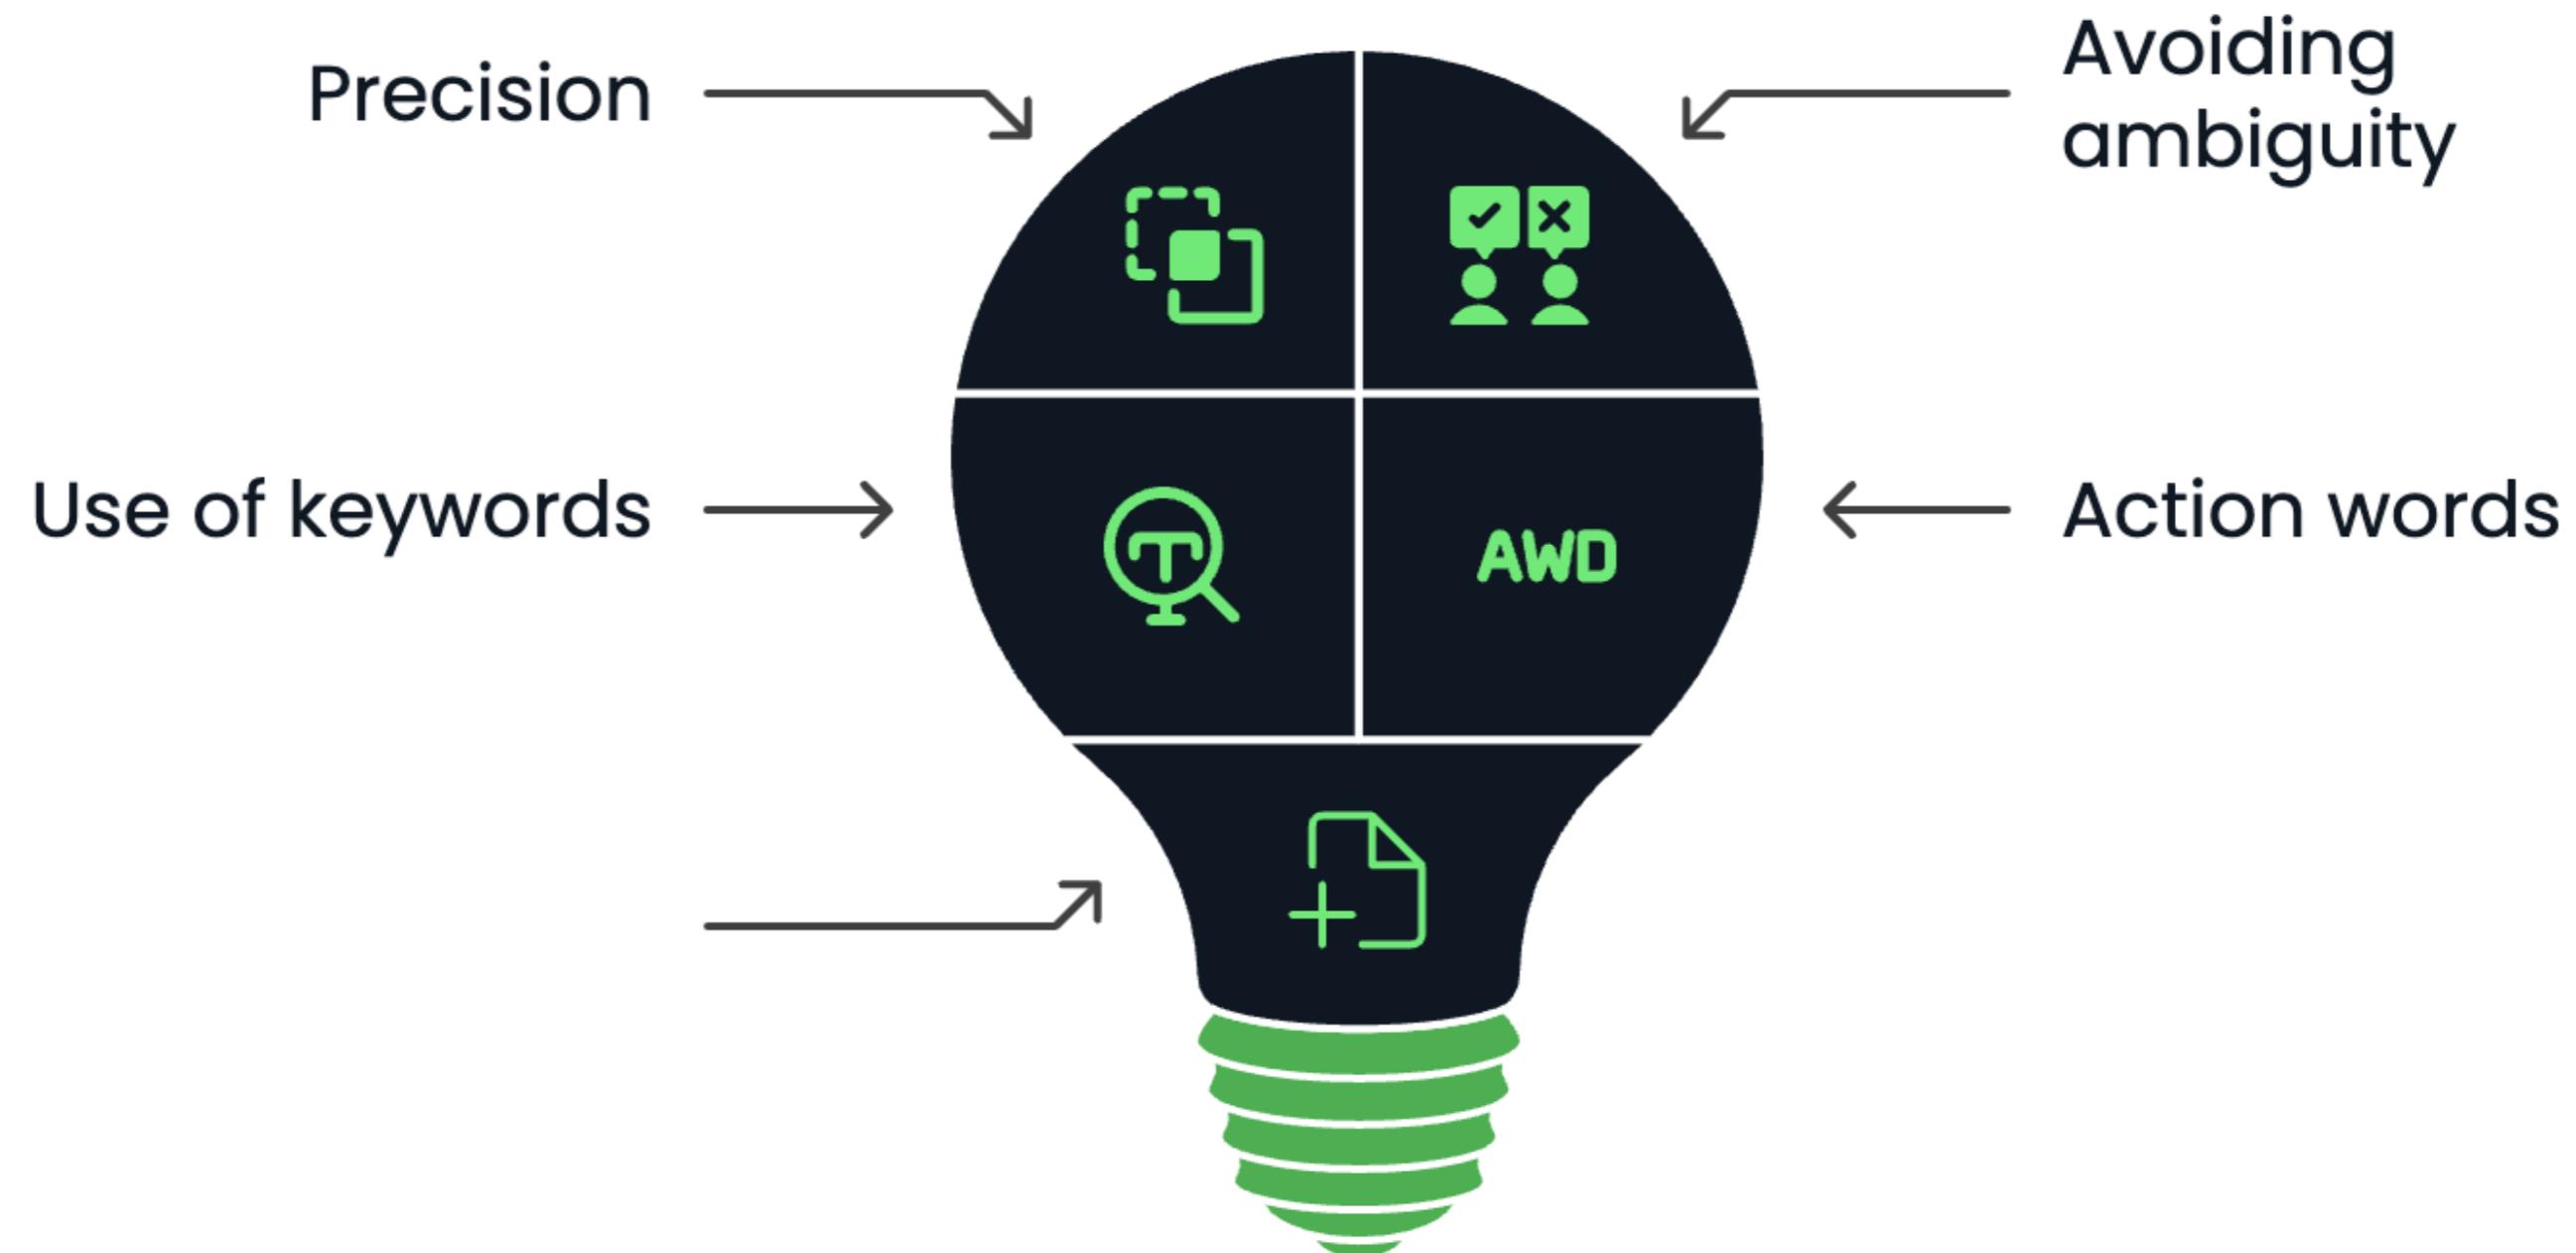
Components of effective prompting



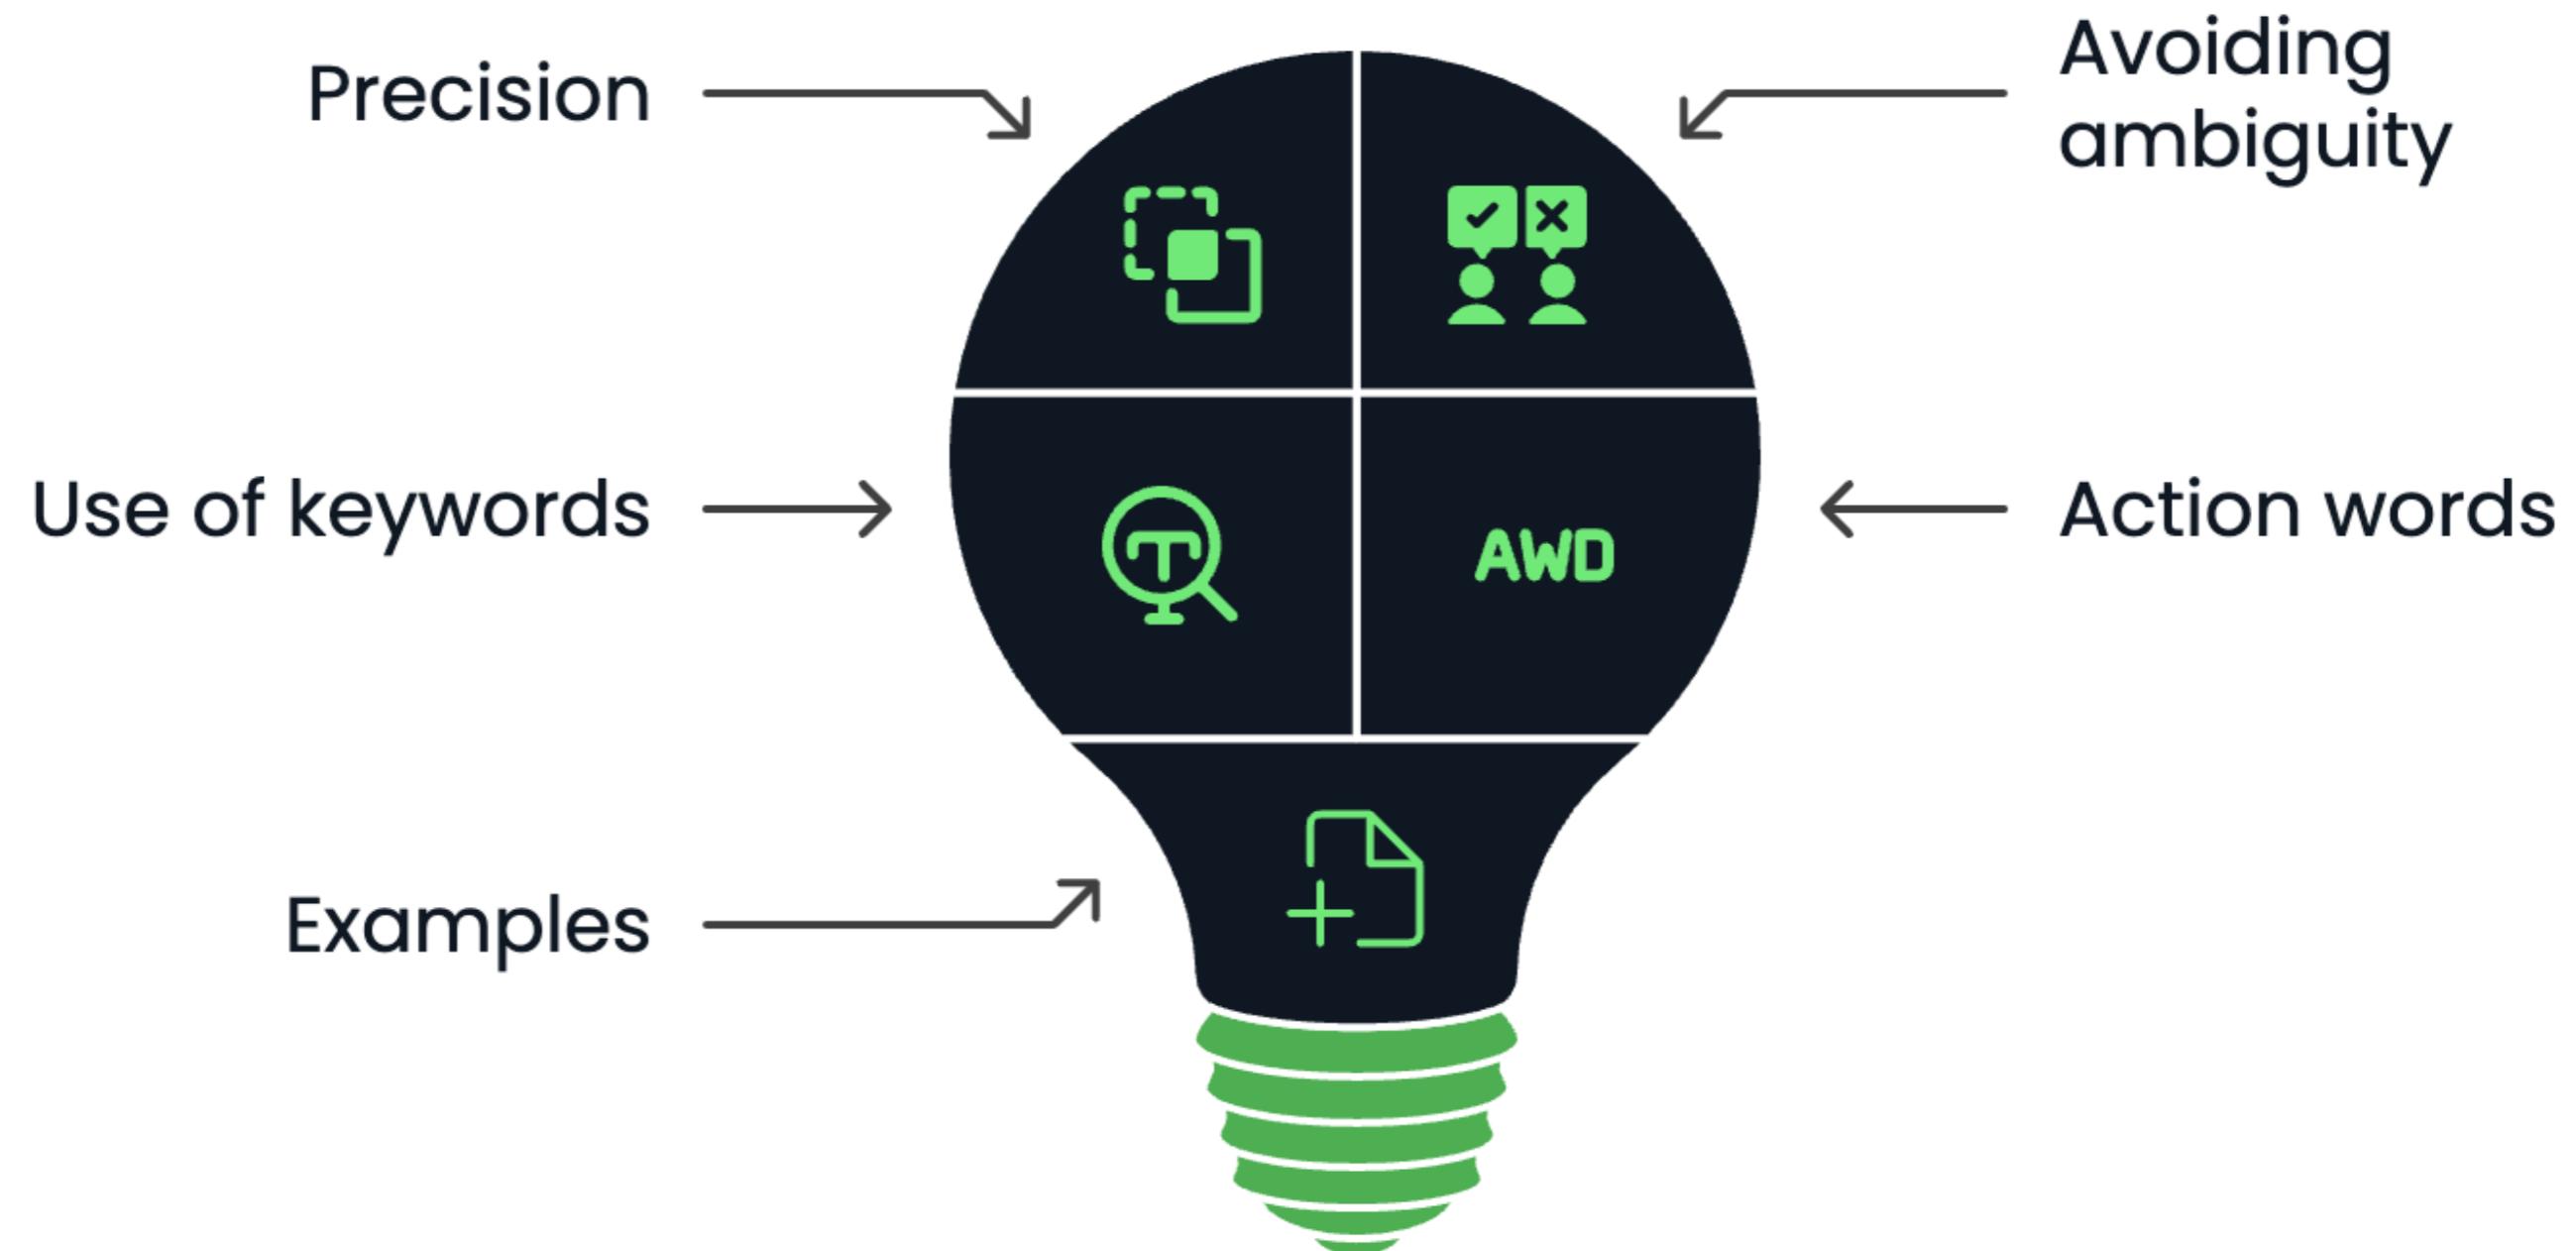
Components of effective prompting



Components of effective prompting



Components of effective prompting



Zero-shot prompting

- Zero-shot prompting: a single instruction

```
text = "Summarize recent mergers in the airline industry."
```

```
output = llm(text)
```

```
print(output['choices'][0]['text'])
```

Recent mergers in the airline industry include Alaska Air Group's acquisition of Hawaiian Airlines in 2024, with both airlines continuing to operate as separate brands. In 2022, Delta Air Lines purchased 20% of LATAM Airlines Group...

Refining zero-shot prompts

- Need to distinguish task, expected output, additional context
- Zero-shot prompting with labels

```
text = """  
    INSTRUCTION: Write concisely and in 2-3 sentences that cover only key points.  
    QUESTION: Summarize recent mergers in the airline industry.  
    ANSWER:  
    """
```

Few-shot prompting

- Few-shot prompting: use of multiple examples

```
text = """  
Aircraft Model: Boeing 787-9  
Passenger Capacity: 296  
Fuel Consumption: 2.5 liters per seat per 100 km
```

```
Aircraft Model: Airbus A321XLR  
Passenger Capacity: 244  
Fuel Consumption: 2.9 liters per seat per 100 km
```

```
Aircraft Model:  
"""
```

Few-shot prompting

- Few-shot prompting: use of multiple examples

```
text = """  
Aircraft Model: Boeing 787-9  
Passenger Capacity: 296  
Fuel Consumption: 2.5 liters per seat per 100 km
```

```
Aircraft Model: Airbus A321XLR  
Passenger Capacity: 244  
Fuel Consumption: 2.9 liters per seat per 100 km
```

```
Aircraft Model:  
"""
```

Few-shot prompting

- Few-shot prompting: use of multiple examples

```
text = """  
Aircraft Model: Boeing 787-9  
Passenger Capacity: 296  
Fuel Consumption: 2.5 liters per seat per 100 km
```

```
Aircraft Model: Airbus A321XLR  
Passenger Capacity: 244  
Fuel Consumption: 2.9 liters per seat per 100 km
```

```
Aircraft Model:  
"""
```

Few-shot prompting

- Few-shot prompting: use of multiple examples

```
text = """  
Aircraft Model: Boeing 787-9  
Passenger Capacity: 296  
Fuel Consumption: 2.5 liters per seat per 100 km
```

```
Aircraft Model: Airbus A321XLR  
Passenger Capacity: 244  
Fuel Consumption: 2.9 liters per seat per 100 km
```

```
Aircraft Model:  
"""
```

Few-shot prompting

```
output = llm(f"Continue the entries: {text}")  
  
print(output['choices'][0]['text'])
```

Aircraft Model: Airbus A350-900

Passenger Capacity: 350

Fuel Consumption: 2.39 liters per seat per 100 km

Using stop words

- Need concise insights
- Use `stop` words to end the response at a specific point
- Example: question-answering application

```
text = "Which airlines operate direct flights from London to Singapore?"
```

```
output = llm(text, stop=["Q:]) # Stop responses at "Q:"  
print(output['choices'][0]['text'])
```

You can fly direct from London to Singapore with Singapore Airlines and British Airways.

Let's practice!

WORKING WITH LLAMA 3

Generating structured output

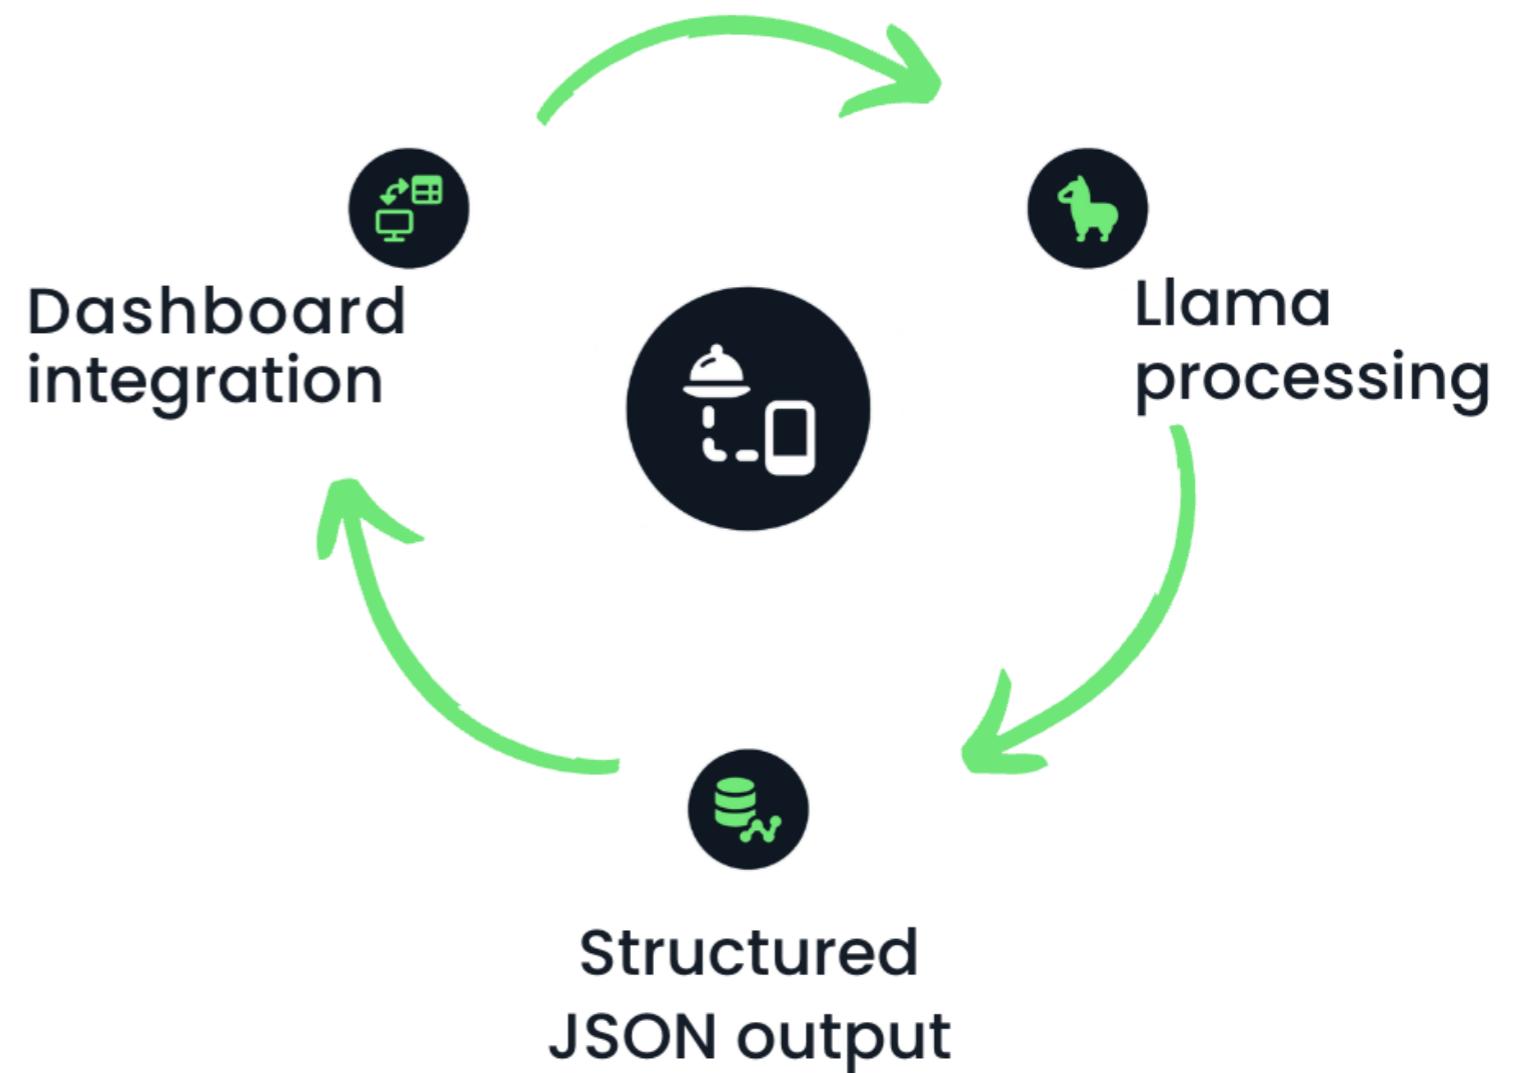
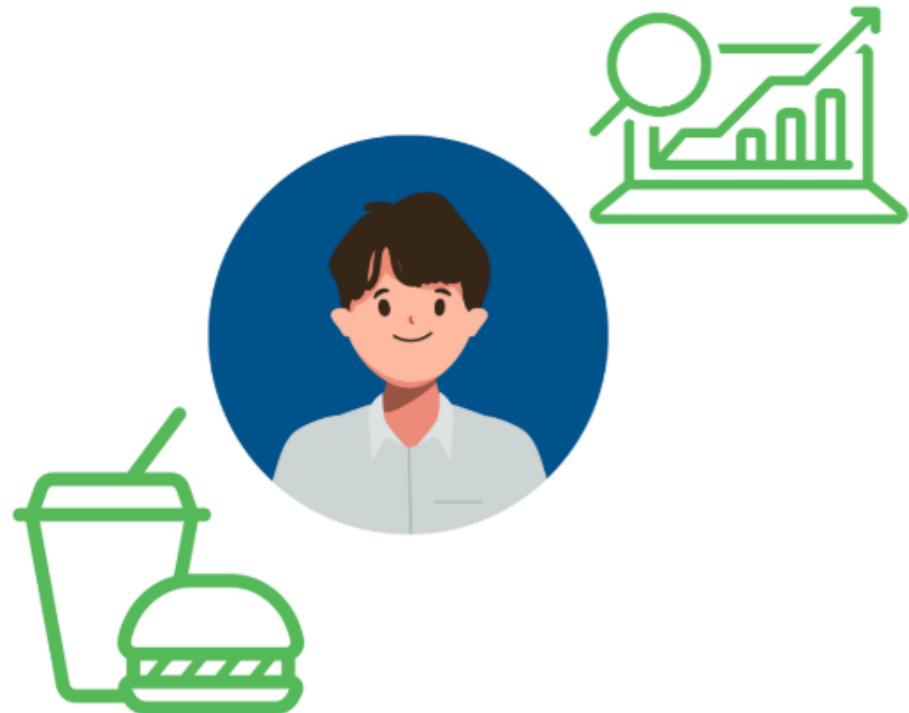
WORKING WITH LLAMA 3



Imtihan Ahmed
Machine Learning Engineer

Structured output in JSON

- Example: Llama responses might be input to a dashboard
- Plain text responses won't work
- We need structured outputs



JSON responses with chat completion

```
response_format = {"type": "json_object"}  
  
message_list = [  
    {"role": "system", # System role defined as market analyst  
        "content": "You are a food industry market analyst. You  
analyze sales data and generate structured JSON reports  
of top-selling beverages."},  
    {"role": "user", # User role to pass the request  
        "content": "Provide a structured report on the top-selling beverages  
this year."}  
]
```

JSON responses with chat completion

```
output = llm.create_chat_completion(  
    messages = message_list,  
    response_format = "json_object"  
)
```

- Response **format** specified as JSON
- Llama generates structured response, no free-flowing text

Extracting the JSON response

```
print(output['choices'][0]['message']['content'])
```

```
{
  "report_name": "Top-Selling Beverages 2024",
  "top_beverages": [
    {
      "rank": 1,
      "beverage_name": "Coca-Cola Classic",
      "sales_volume": 2.1,
      "growth_rate": 1.9
    }, ...
  }
}
```

Defining a schema

```
response_format = {  
    "type": "json_object",  
    "schema": {  
        "type": "object",  
        "properties": {  
            "Product Name": {"type": "string"},  
            "Category": {"type": "string"},  
            "Sales Growth": {"type": "float"}  
        }  
    }  
}
```

- Can specify a **schema**: rules to define how the data should be formatted

Defining a schema

```
output = llm.create_chat_completion(  
    messages = message_list,  
    response_format = response_format)  
  
print(output['choices'][0]['message']['content'])
```

```
{  
    "Product Name": "Coca-Cola",  
    "Category": "Soft Drink",  
    "Sales Growth": 12.5  
}
```

Let's practice!

WORKING WITH LLAMA 3

Building conversations

WORKING WITH LLAMA 3



Imtihan Ahmed
Machine Learning Engineer

Maintaining context



User inquiry

Maintaining context



User inquiry

AI response

Maintaining context



User inquiry

AI response

User follow-up

Maintaining context



User inquiry



AI response



User follow-up



AI memory
use

Maintaining context



User inquiry



AI response



User follow-up



AI memory
use



AI response

- Track a chat history with a `Conversation` class

Conversation class

- Can store a history of prior messages

```
class Conversation:  
    def __init__(self, llm: Llama, system_prompt='', history[]):  
        self.llm = llm  
        self.system_prompt = system_prompt  
        self.history = [{"role": "system", "content": self.system_prompt}] + history
```

Conversation class

- Can store a history of prior messages

```
class Conversation:  
    def __init__(self, llm: Llama, system_prompt='', history[]):  
        self.llm = llm  
        self.system_prompt = system_prompt  
        self.history = [{"role": "system", "content": self.system_prompt}] + history  
  
    def create_completion(self, user_prompt ''):  
        self.history.append({"role": "user", "content": user_prompt}) # Append input  
        output = self.llm.create_chat_completion(messages=self.history)  
        conversation_result = output['choices'][0]['message']  
        self.history.append(conversation_result) # Append output  
        return conversation_result['content'] # Return output
```

Running a multi-turn conversation

```
conversation = Conversation(llm, system_prompt="You are a virtual travel assistant  
helping with planning trips.")  
  
response1 = conversation.create_completion("What are some destinations in France for a  
short weekend break?")  
print(f"Response 1: {response1}")  
  
response2 = conversation.create_completion("How about Spain?")  
  
print(f"Response 2: {response2}")
```

Running a multi-turn conversation

```
print(f"Response 1: {response1}")
```

```
print(f"Response 2: {response2}")
```

Response 1: France is ideal for a short weekend break:

1. **Paris**: The capital city is famous for its iconic landmarks like the ...
2. **Provence**: Known for its beautiful landscapes, mild climate, and ...

Response 2: Here are some destinations in Spain for a short weekend break:

1. **Barcelona**: Visit the famous landmarks like the Sagrada Familia, ...

Running a multi-turn conversation

```
print(f"Response 1: {response1}")
```

```
print(f"Response 2: {response2}")
```

Response 1: France is ideal for a short **weekend break**:

1. **Paris**: The capital city is famous for its iconic landmarks like the ...
2. **Provence**: Known for its beautiful landscapes, mild climate, and ...

Response 2: Here are some destinations in Spain for a short **weekend break**:

1. **Barcelona**: Visit the famous landmarks like the Sagrada Familia, ...

Let's practice!

WORKING WITH LLAMA 3

Congratulations!

WORKING WITH LLAMA 3



Imtihan Ahmed
Machine Learning Engineer

Let's recall

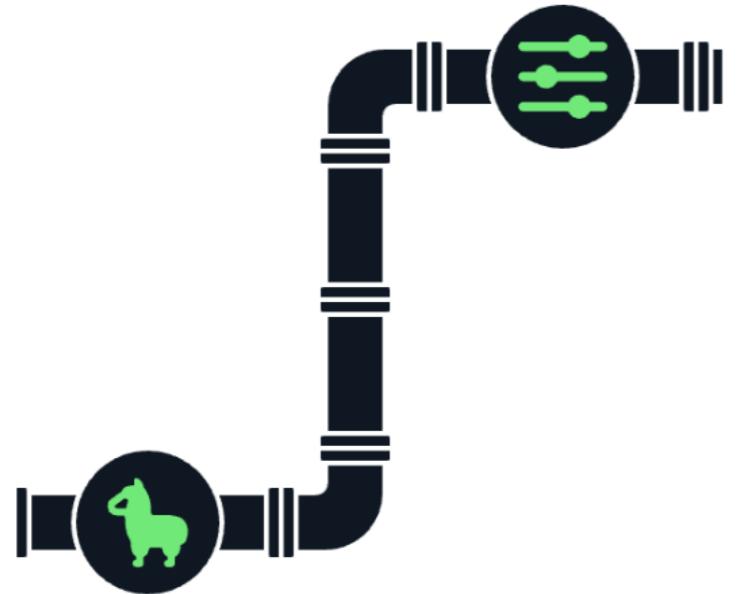


01

Run Llama
locally

```
from llama_cpp import Llama  
  
llm = Llama(model_path = "path/to/model.gguf")
```

Let's recall



01

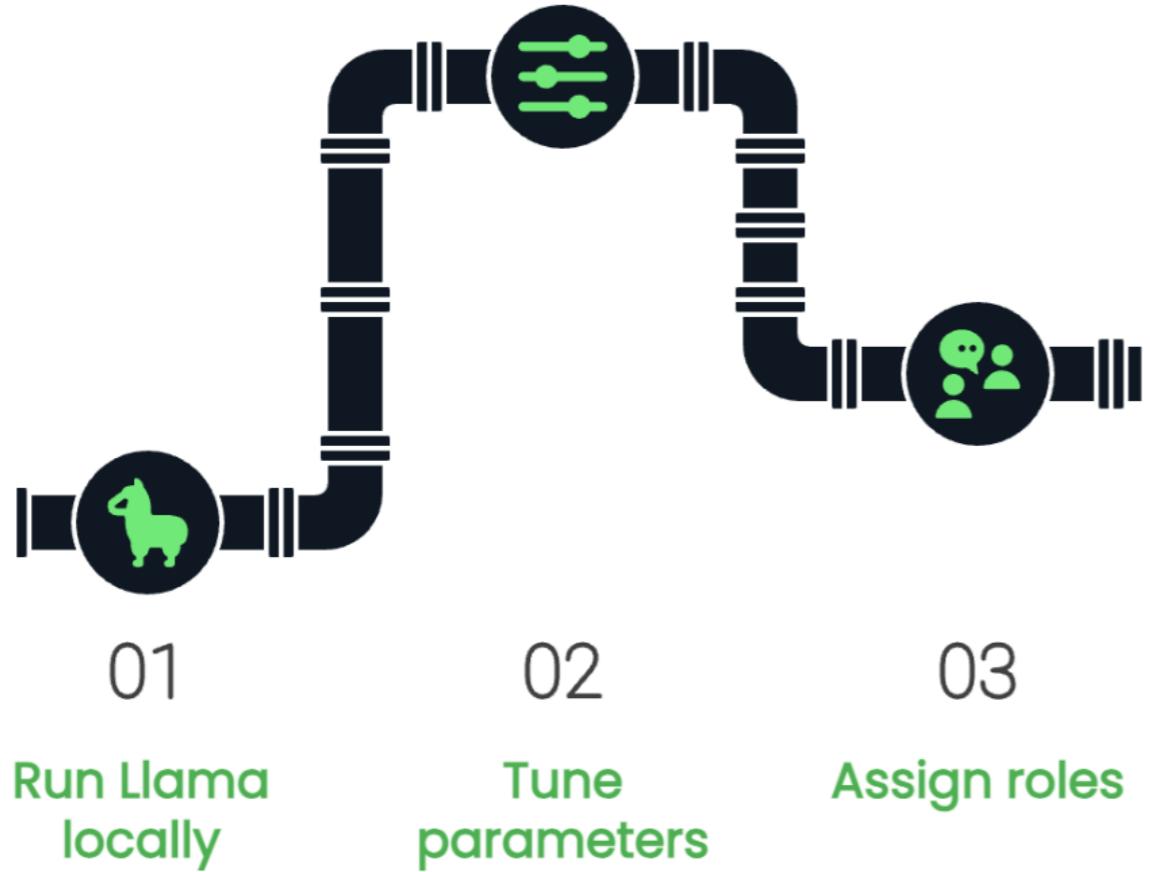
Run Llama
locally

02

Tune
parameters

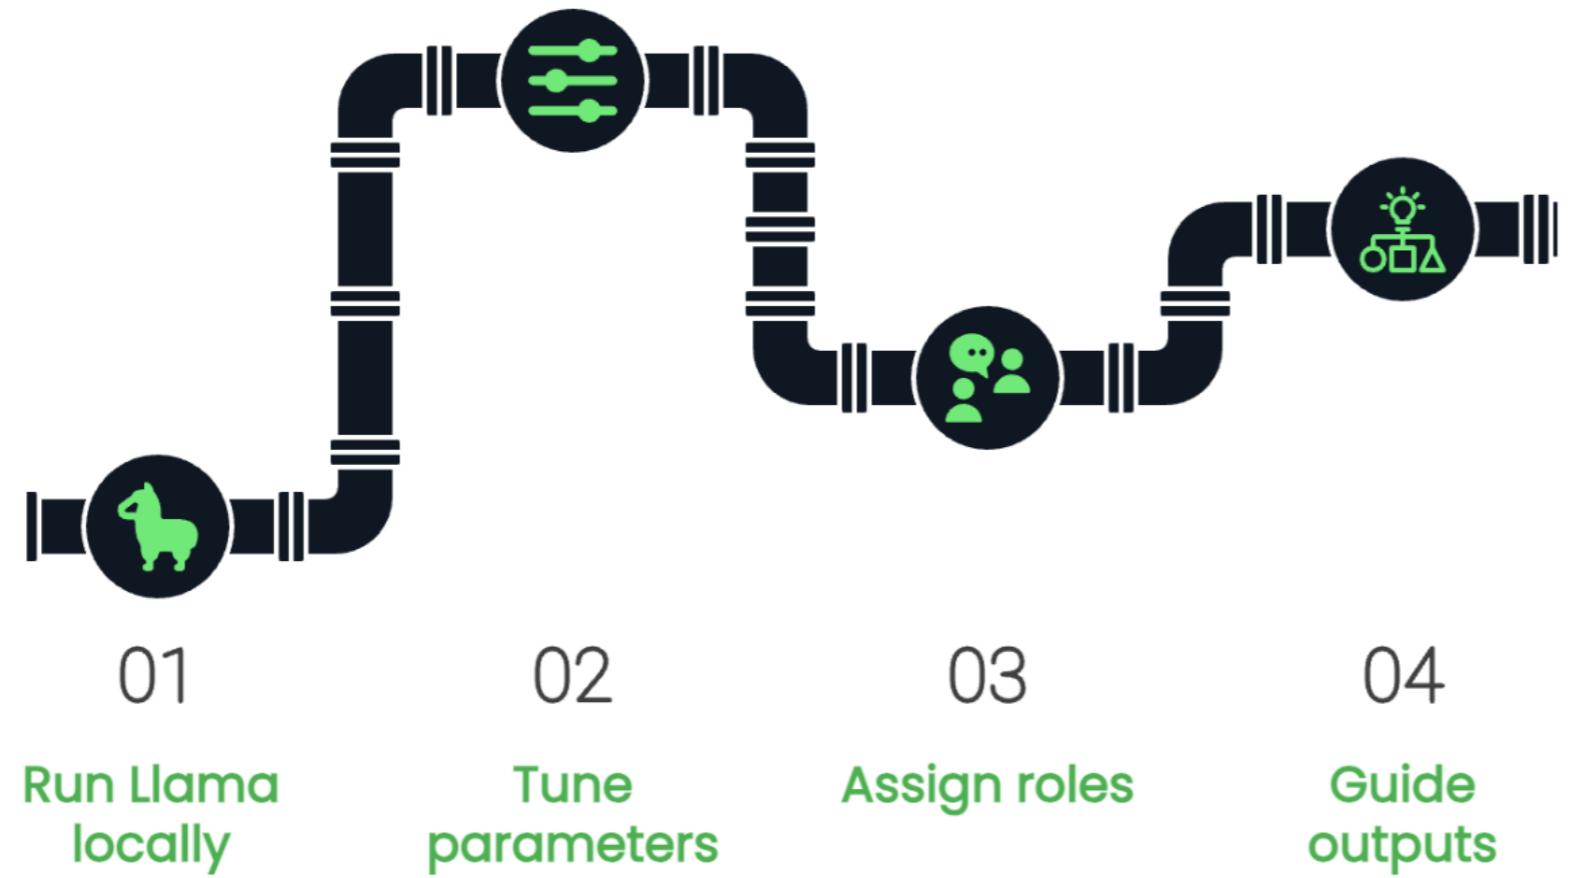
- `temperature` , `top_k` , `top_p` parameters

Let's recall



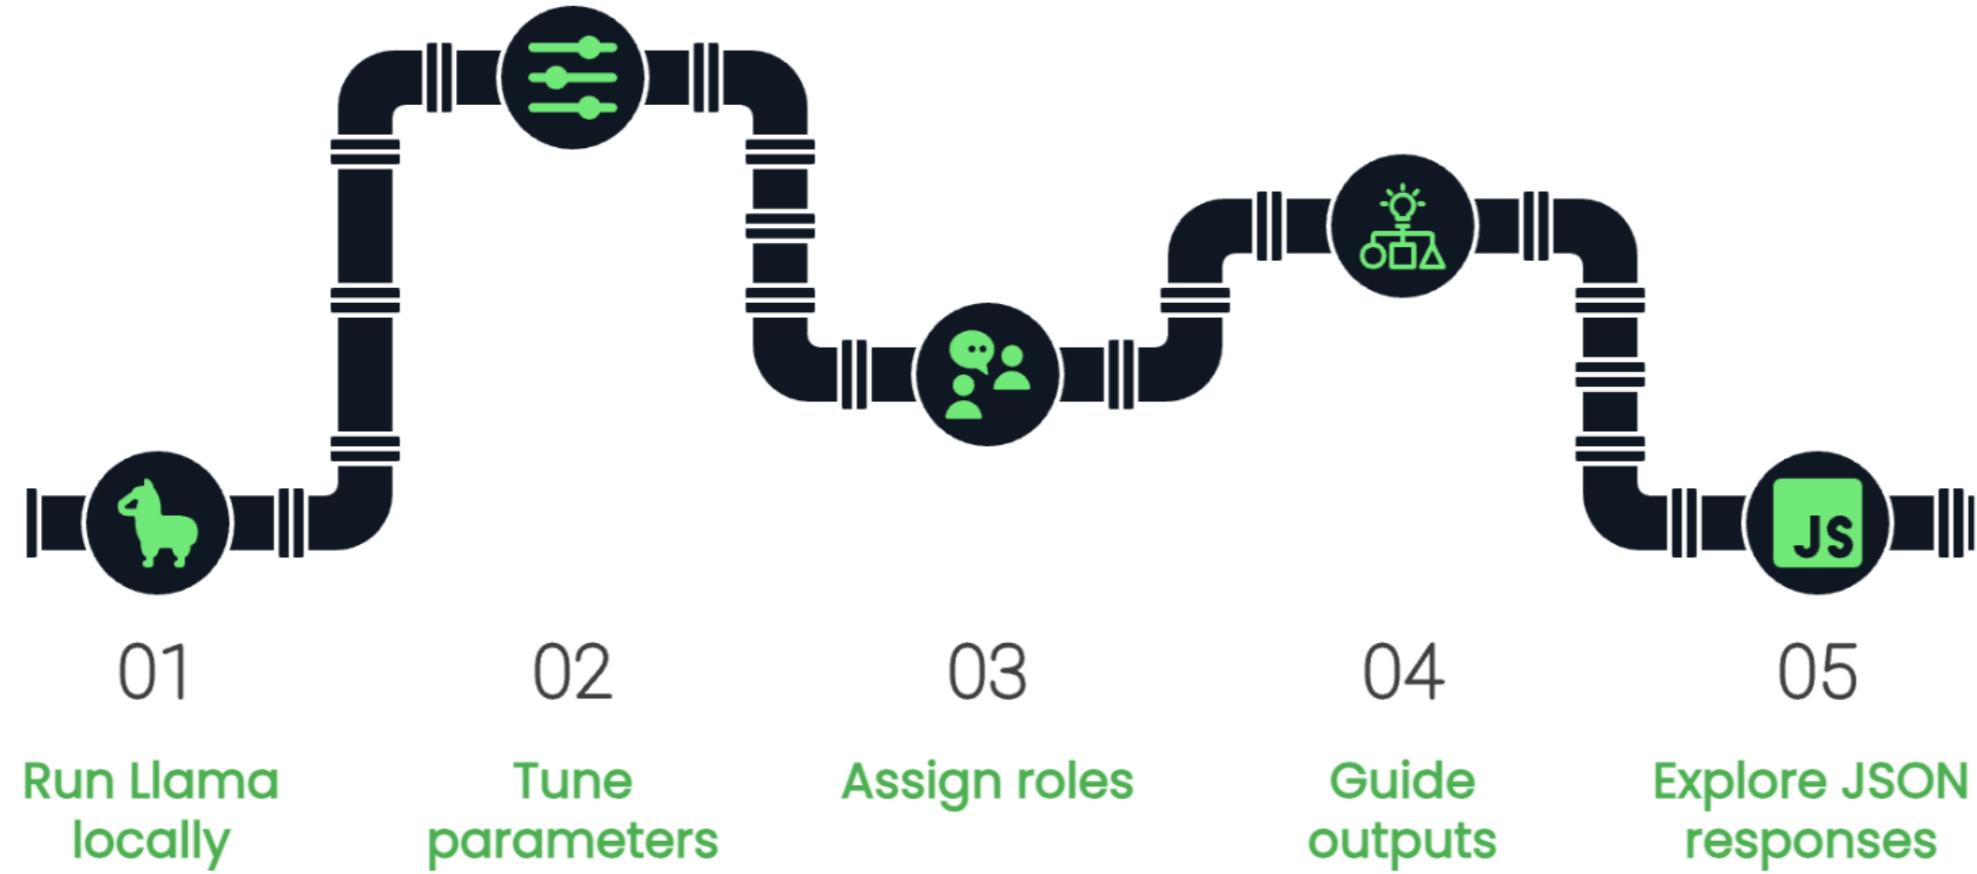
```
message_list = [{"role": "system", "content": system_message},  
                 {"role": "user", "content": user_message}]
```

Let's recall



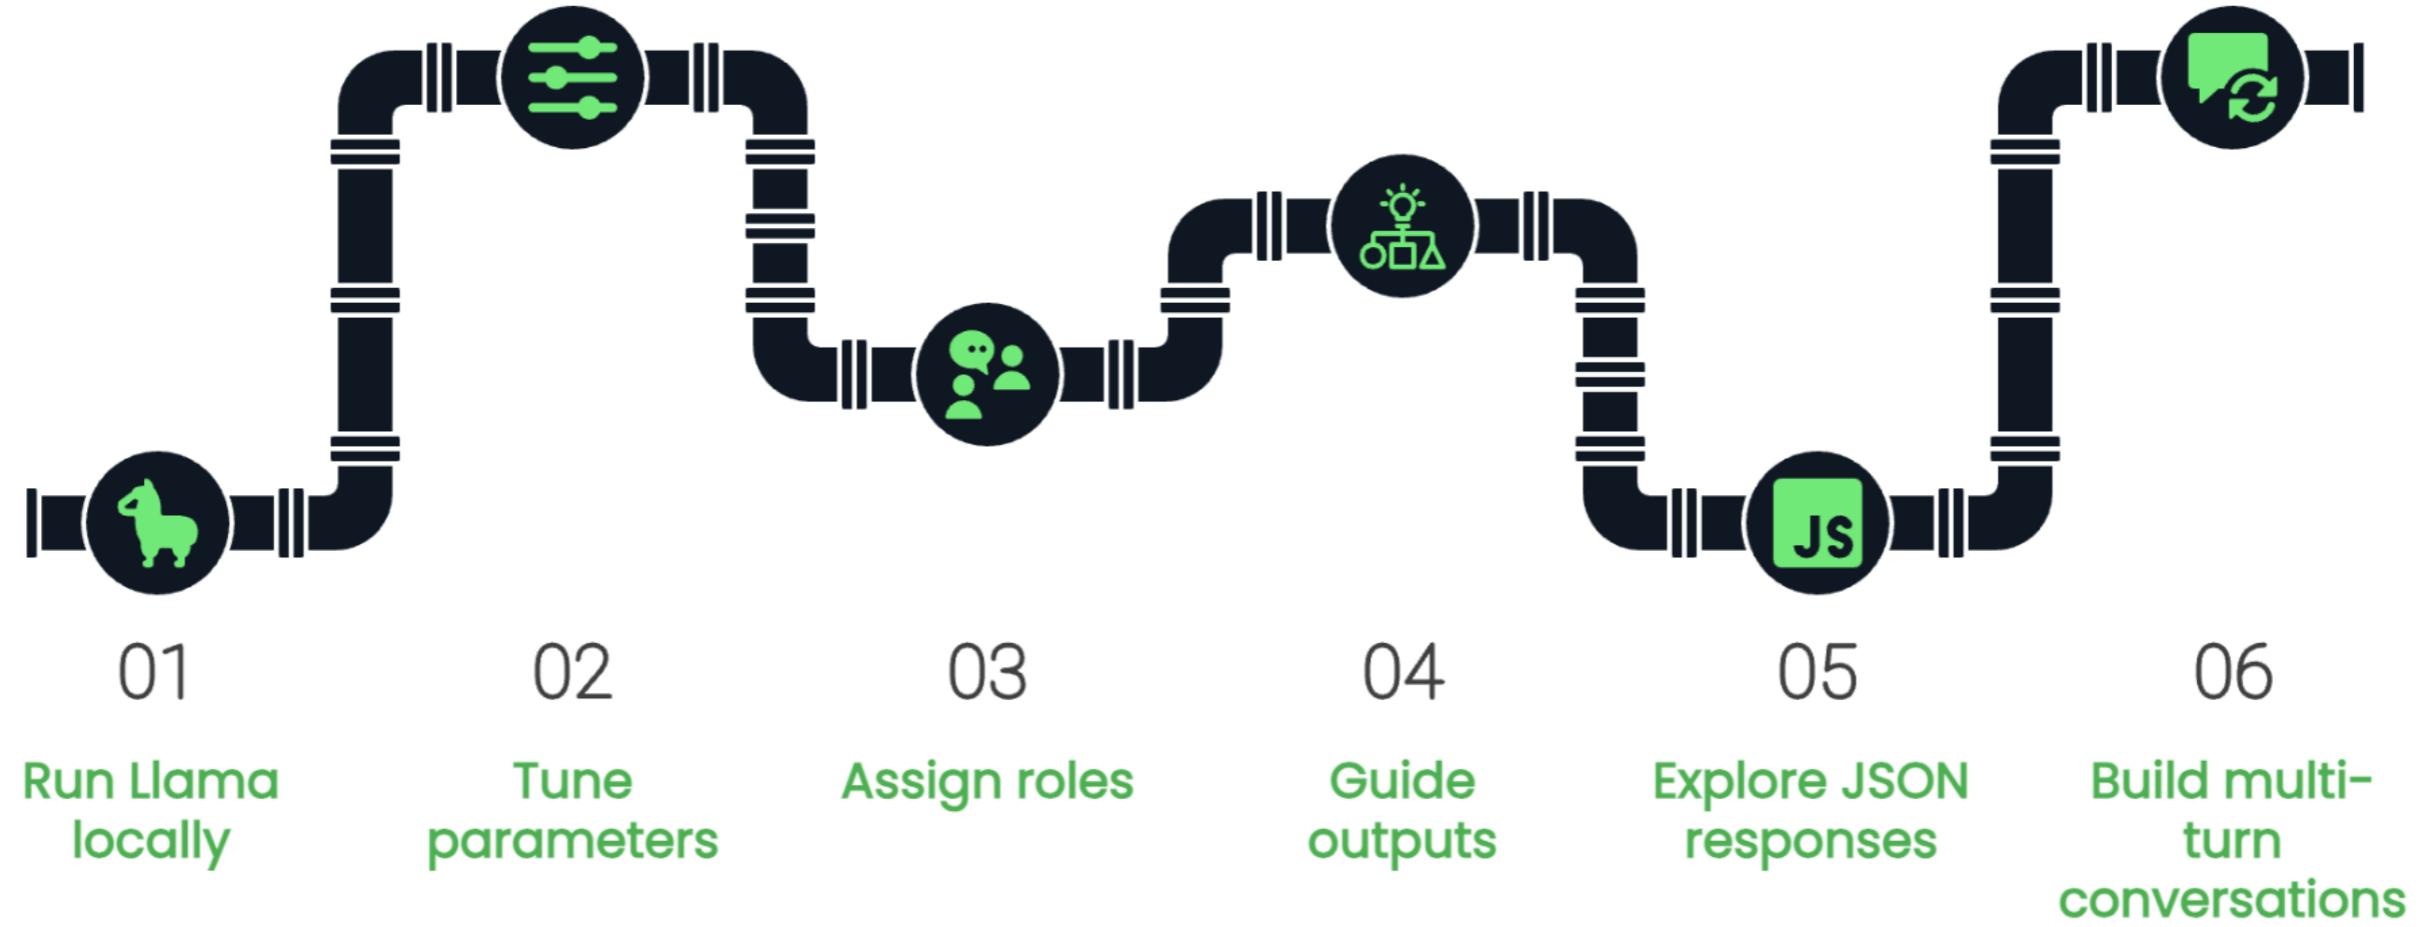
- Precise prompts
- stop words
- Zero-shot/Few-shot prompting

Let's recall



```
response_format = {"type": "json_object"}
```

Let's recall



- Conversation class
- .create_completion() method

What's next?

SKILL TRACK

Llama Fundamentals

[Enroll in Track](#)

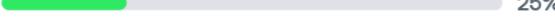
∞ Llama ⏱ 5 hours ⚡ 2 courses 📂 1 project 📃 588 participants

Track Description

Experiment with Llama 3 to run inference on pre-trained models, fine-tune them on custom datasets, and optimize performance.

COURSE

1 Working with Llama 3



25%

PROJECT

Φ BONUS Classifying Emails using Llama

Go further! Gain mastery with this optional material.

COURSE

3 Fine-Tuning with Llama 3

INSTRUCTORS

Imtihan Ahmed
Machine Learning Engineer

Francesca Donadoni
AI Curriculum Manager at DataCamp

PREMIUM PROJECT

Classifying Emails using Llama

Build an AI-powered inbox assistant to classify your emails using Llama.

[Start Project](#)

∞ 1 tasks 📃 287 participants 1,500 XP

Project Description

Build an intelligent email assistant using Llama's local LLM capabilities. Using prompt engineering and model integration, develop a system to classify emails into predefined categories to help you prioritize tasks and manage your inbox more efficiently.

PREREQUISITES

Working with Llama 3

INSTRUCTORS

Dheeraj Agrawal
Instructor

[See All](#)

Llama Fundamentals

Classifying Emails using Llama

Thank you!

WORKING WITH LLAMA 3